

T5ynth macOS Installation

Short guide for the current unsigned macOS build with the usual Gatekeeper override step.

The screenshots show the real installation flow on a Mac with a German user interface.

Unfortunately I only had German screenshots available for this documentation pass, but the sequence is identical on English-language systems.

Tip: the recommended default model, Stable Audio 3, downloads directly in-app – see Appendix C. Steps 11–15 instead show the manual installation of Stable Audio Open Small.

Quick summary

- Open the .pkg
- If blocked: Privacy & Security > Open Anyway
- Click through the installer normally
- Launch T5ynth.app from /Applications

Step 1: Open the package

If macOS blocks the package the first time, that is expected for this current build.



Step 2: Open Privacy & Security

In System Settings > Privacy & Security, scroll down and click “Open Anyway” for T5ynth.

„T5ynth-mac...2-test.pkg“ wurde blockiert, um deinen Mac zu schützen.

Dennoch öffnen

Apple konnte nicht überprüfen, ob „T5ynth-mac...2-test.pkg“ frei von Schadsoftware ist, die deinen Mac beschädigen oder deine Privatsphäre verletzen kann.

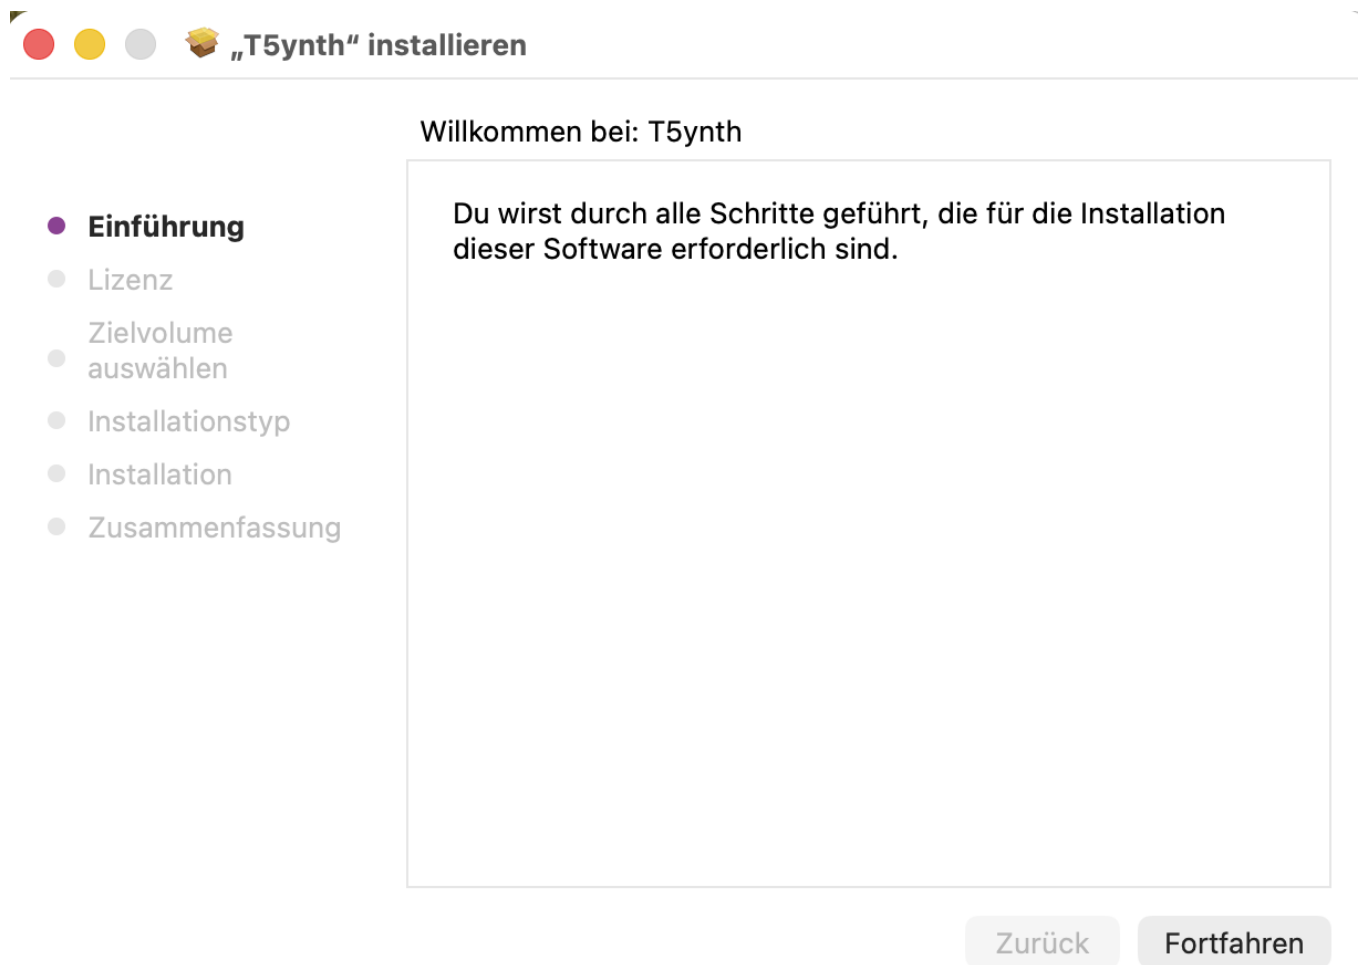
Step 3: Confirm opening

Confirm the security prompt once more with "Open Anyway".



Step 4: The installer starts normally

After that, the normal macOS installer opens.



Step 5: License page

Continue to the license page.



Softwarelizenzvertrag

- Einführung
- **Lizenz**
- Zielvolume auswählen
- Installationstyp
- Installation
- Zusammenfassung

T5ynth — Text-to-Sound Synthesizer
Copyright (C) 2024-2026 Prof. Dr. Benjamin Jörisen / AI4ArtsEd
(UCDCAE AI Lab)
<https://github.com/joeriben/t5ynth>

T5ynth is a novel synthesizer concept that navigates the T5 text embedding space of a diffusion audio model, enabling direct sonic exploration of semantic dimensions. Conceived and developed by Prof. Dr. Benjamin Jörisen as part of the AI4ArtsEd project at the UCDCAE AI Lab.

This program is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

Drucken ...

Sichern ...

Zurück

Fortfahren

Step 6: Accept the license

Without “Accept”, the installation cannot continue.

Um mit der Installation der Software fortzufahren, musst du den Softwarelizenzvertrag akzeptieren.

Klicke auf „Akzeptieren“, um mit der Installation fortzufahren.

Klicke auf „Ablehnen“, um die Installation abubrechen und zu beenden.


Lizenz lesen

Ablehnen

Akzeptieren

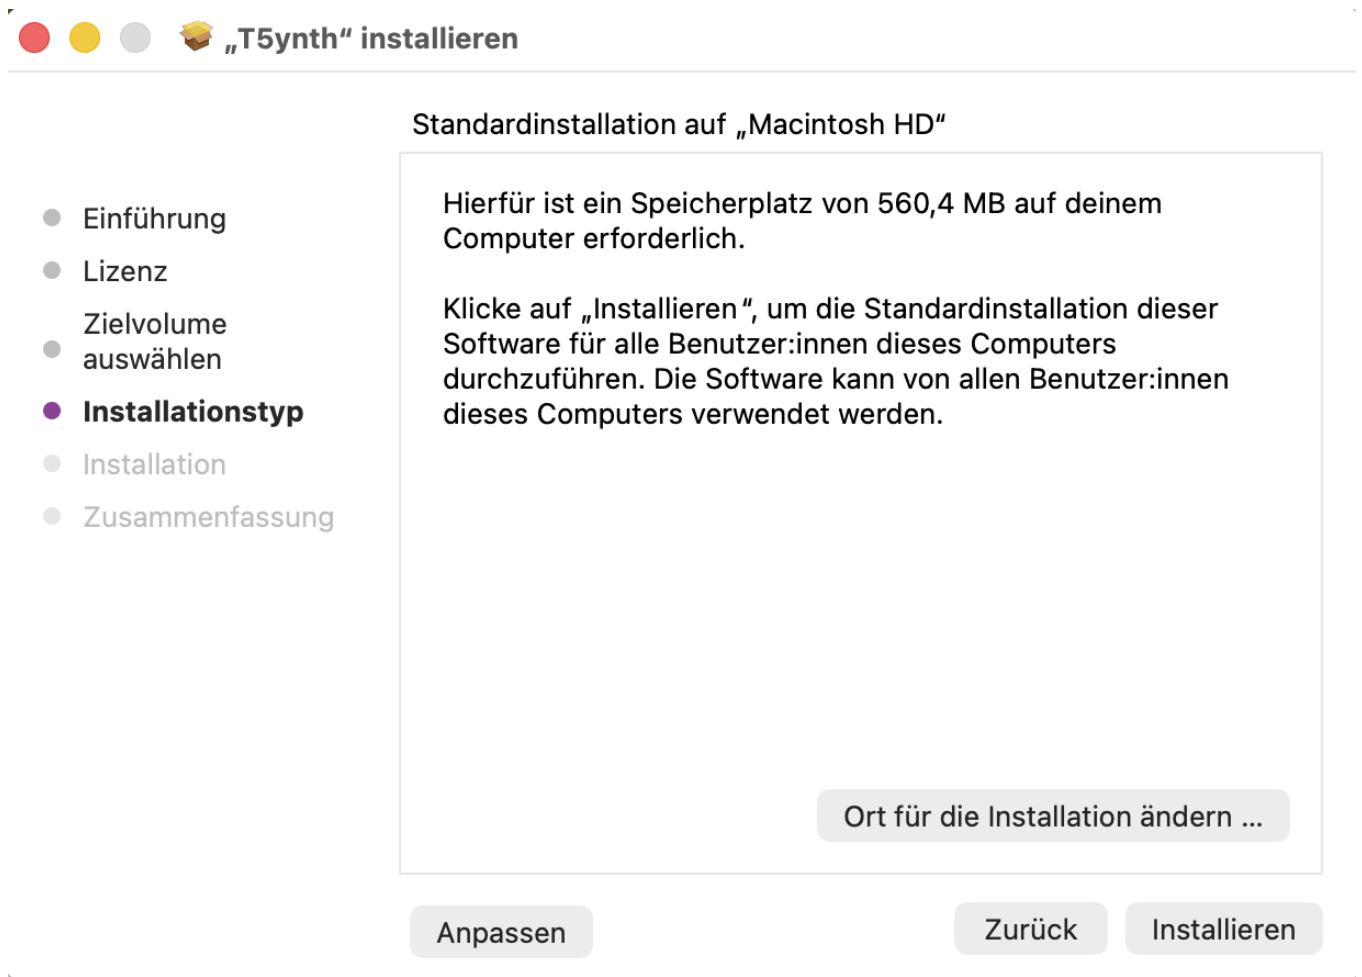
Step 7: Choose the install target

In the standard case, choose installation for all users of this computer.



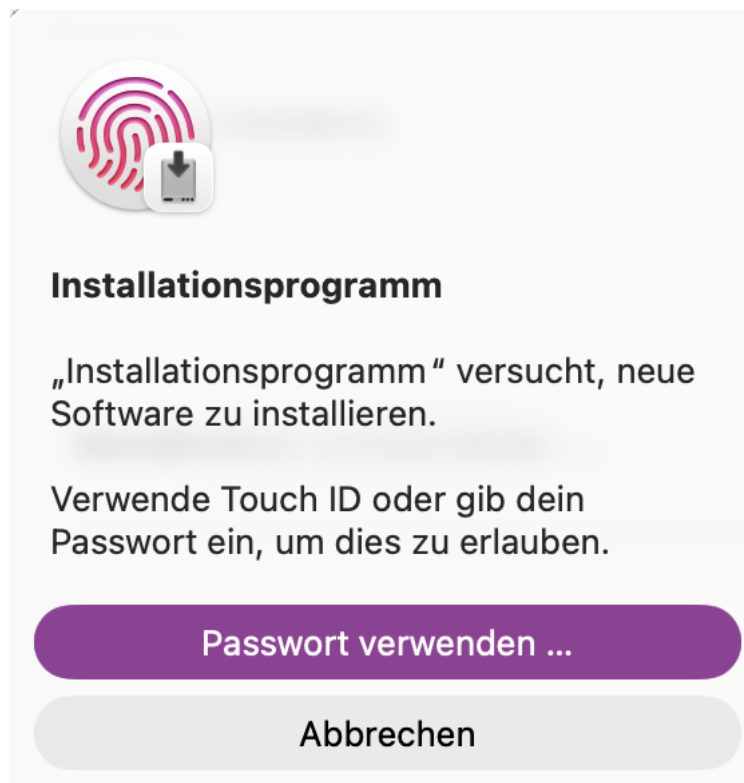
Step 8: Start the installation

Click "Installieren" to begin the standard installation on the system volume.




Step 9: Confirm with password or Touch ID

macOS asks for administrator permission for the actual installation.




Step 10: Installation complete

Once this dialog appears, T5ynth.app is correctly installed in /Applications.




Step 11: Open T5ynth and check the model state

In the Model Manager, Stable Audio Open Small is initially still not installed.




Step 12: Allow access to Downloads

If Auto-Scan should look for the model files in Downloads, allow access to that folder.




Step 13: Import the model files

T5ynth copies the model files from Downloads into its internal model directory.




Step 14: Backend starts

After the copy step, T5ynth briefly shows the transition state “Backend: Starting...”.



Step 15: Model is ready

Once “Installed” and “Backend: Connected” appear, T5ynth is ready to use.



Appendix C: Loading Stable Audio 3 (recommended default)

Stable Audio 3 is T5ynth's current default model (version 2.1.0). There are two variants: Small Music for instrumental music and Small SFX for sound effects; both share one architecture and one t5gemma text encoder.

Before downloading, T5ynth shows two licenses that both must be accepted: the Stability AI Community License for the audio model and the Google Gemma Terms of Use plus Prohibited Use Policy for the t5gemma encoder. The weights then download without a HuggingFace account or token.

- Select “Stable Audio 3 Small Music” or “Small SFX” in the Model Manager.
- Accept both licenses with “Accept & Download”; the download starts.
- Both variants share the same t5gemma text encoder.
- Status switches to “Installed”, then “Backend: Connected”.

Note

Stable Audio 3 reads up to 256 tokens per prompt. T5ynth automatically prepends the required modality prefix (“TrackType: Music, ...” or “TrackType: SFX, ...”). SA3 was trained on comma-joined, field-tagged metadata, so structured prompts such as “Instruments: synth pad, Moods: warm” fit its training distribution; plain descriptions work too.

Appendix B: Loading Stable Audio Open 1.0

Stable Audio Open 1.0 is the original full-size Stable Audio Open checkpoint. It installs directly inside T5ynth, like Stable Audio 3.

Before downloading, T5ynth shows the Stability AI Community License, which must be accepted. The ~4.85 GB checkpoint then downloads without a HuggingFace account or token, and a copy of the license is written into the model folder.


- Select “Stable Audio Open 1.0” in the Model Manager.
- Accept the license with “Accept & Download”; the download starts.
- Status switches to “Installed”, then “Backend: Connected”.
- Already have the files from Stability by hand? Use “Auto-Scan” instead.

Note

Like Stable Audio Open Small, Stable Audio Open 1.0 also needs the separate “T5-Base text encoder” (ungated, Apache-2.0) – install it from the Model Manager too, or the backend will not load. Stable Audio Open 1.0 is also much larger than Small; for first tests on a new Mac, Stable Audio 3 or Stable Audio Open Small is usually the quicker starting point.


Appendix A1: Select AudioLDM2

In the Model Manager, choose AudioLDM2. Unlike Stable Audio Open Small, T5ynth can download this model directly from HuggingFace.




Appendix A2: Confirm the license

Before downloading, T5ynth shows the AudioLDM2 license terms. Click “Accept & Download” to continue.



Appendix A3: Download in progress

While the download is running, T5ynth shows the current file name and progress.



The screenshot shows the T5ynth Model Manager interface. At the top, there is a 'Model Manager' section with a dropdown menu set to 'AudioLDM2'. Below this, a status message indicates 'audioldm2: Not installed'. A green status bar shows 'Backend: Connected'. Below the status bar are three buttons: 'Auto-Scan', 'Browse...', and 'Open Model Page'. A large green button labeled 'Download from HuggingFace' is prominently displayed. Below this button, a progress bar shows 'Downloading: language_model/model.safetensors (3/26)'. The main content area displays the model name 'AUDIOLDM2' and a detailed description: 'Academic latent-diffusion text-to-audio model published by CVSSP / University of Surrey and collaborators (Liu et al., 2023), released as an open research artefact for studying generalised audio, music and speech generation from text. Ungated on HuggingFace and the only engine T5ynth can install directly. Click 'Download from HuggingFace' above and wait for the download to finish.' Below the description, the source URL 'https://huggingface.co/cvssp/audioldm2' and the target path '/Users/vera/Library/T5ynth/models/audioldm2' are listed. At the bottom, the license 'CC BY-NC-SA 4.0 -- non-commercial use only, no revenue threshold, no exceptions.' is displayed.

Model Manager AudioLDM2

audioldm2: Not installed

Backend: Connected

Auto-Scan Browse... Open Model Page

Download from HuggingFace

Downloading: language_model/model.safetensors (3/26)


AUDIOLDM2
Academic latent-diffusion text-to-audio model published by CVSSP / University of Surrey and collaborators (Liu et al., 2023), released as an open research artefact for studying generalised audio, music and speech generation from text. Ungated on HuggingFace and the only engine T5ynth can install directly. Click 'Download from HuggingFace' above and wait for the download to finish.

Source: <https://huggingface.co/cvssp/audioldm2>
Target: /Users/vera/Library/T5ynth/models/audioldm2

License: CC BY-NC-SA 4.0 -- non-commercial use only, no revenue threshold, no exceptions.

Appendix A4: Backend starts

After the download, the model is activated automatically and T5ynth briefly shows “Backend: Starting...”.



Appendix A5: AudioLDM2 is ready

Once “Installed” and “Backend: Connected” appear, AudioLDM2 is ready to use.

