

T5ynth macOS Installation

Short guide for the current unsigned macOS build with the usual Gatekeeper override step.

The screenshots show the real installation flow on a Mac with a German user interface.

Unfortunately I only had German screenshots available for this documentation pass, but the sequence is identical on English-language systems.

Tip: every model loads directly in the app, with no HuggingFace account and no token – see the appendices. The recommended default model is Stable Audio 3 (Appendix C).

Quick summary

- Open the .pkg
- If blocked: Privacy & Security > Open Anyway
- Click through the installer normally
- Launch T5ynth.app from /Applications

Step 1: Open the package

If macOS blocks the package the first time, that is expected for this current build.



Step 2: Open Privacy & Security

In System Settings > Privacy & Security, scroll down and click “Open Anyway” for T5ynth.

„T5ynth-mac...2-test.pkg“ wurde blockiert, um deinen Mac zu schützen.

Dennoch öffnen

Apple konnte nicht überprüfen, ob „T5ynth-mac...2-test.pkg“ frei von Schadsoftware ist, die deinen Mac beschädigen oder deine Privatsphäre verletzen kann.

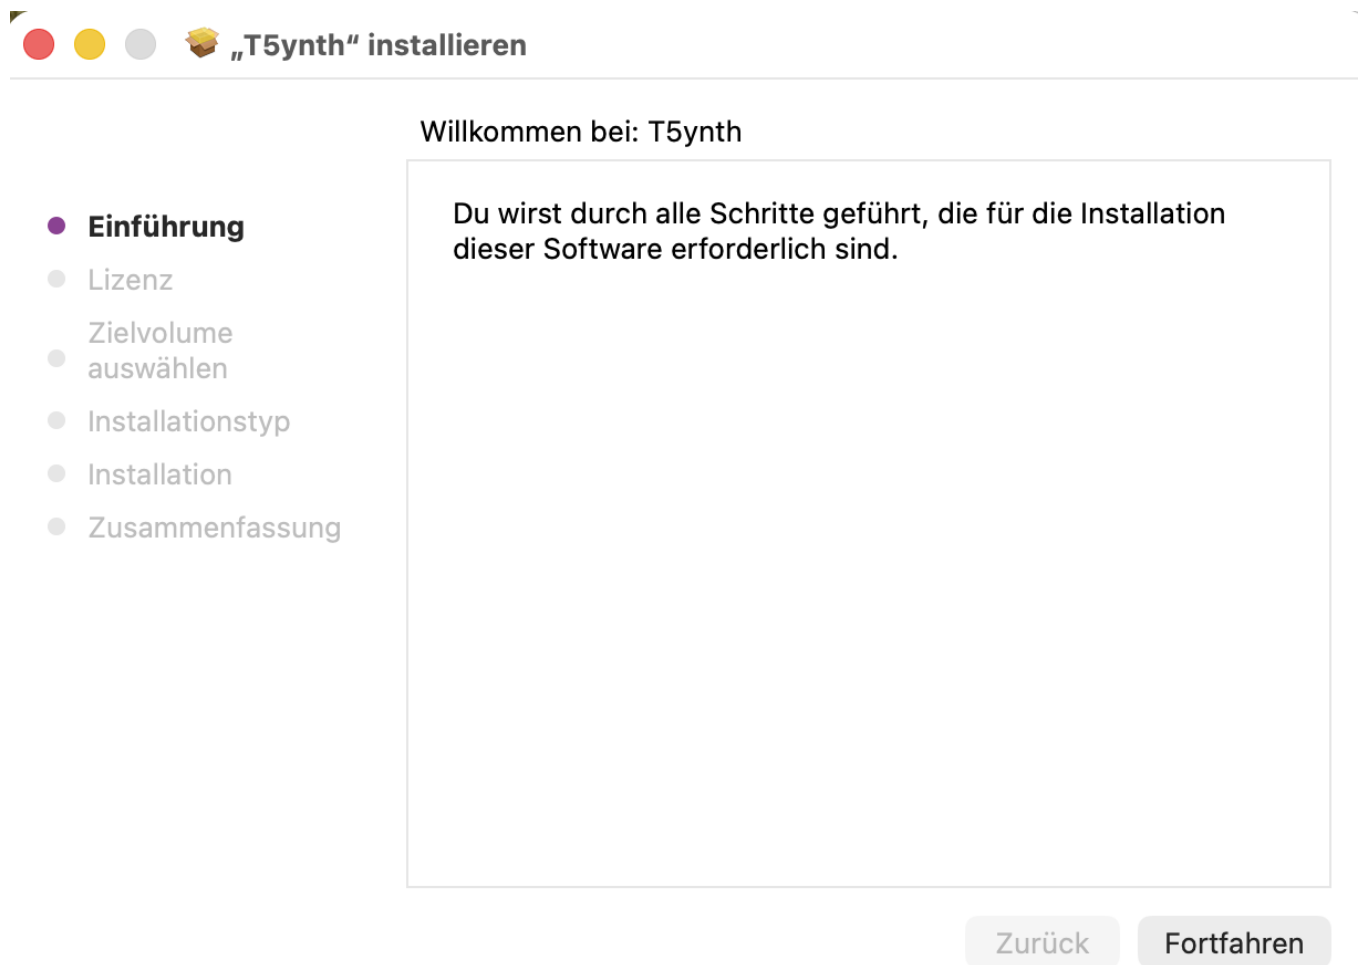
Step 3: Confirm opening

Confirm the security prompt once more with "Open Anyway".



Step 4: The installer starts normally

After that, the normal macOS installer opens.



Step 5: License page

Continue to the license page.



Softwarelizenzvertrag

- Einführung
- **Lizenz**
- Zielvolume auswählen
- Installationstyp
- Installation
- Zusammenfassung

T5ynth — Text-to-Sound Synthesizer
Copyright (C) 2024-2026 Prof. Dr. Benjamin Jörisen / AI4ArtsEd
(UCDCAE AI Lab)
<https://github.com/joeriben/t5ynth>

T5ynth is a novel synthesizer concept that navigates the T5 text embedding space of a diffusion audio model, enabling direct sonic exploration of semantic dimensions. Conceived and developed by Prof. Dr. Benjamin Jörisen as part of the AI4ArtsEd project at the UCDCAE AI Lab.

This program is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

Drucken ...

Sichern ...

Zurück

Fortfahren

Step 6: Accept the license

Without “Accept”, the installation cannot continue.

Um mit der Installation der Software fortzufahren, musst du den Softwarelizenzvertrag akzeptieren.

Klicke auf „Akzeptieren“, um mit der Installation fortzufahren.

Klicke auf „Ablehnen“, um die Installation abubrechen und zu beenden.


Lizenz lesen

Ablehnen

Akzeptieren

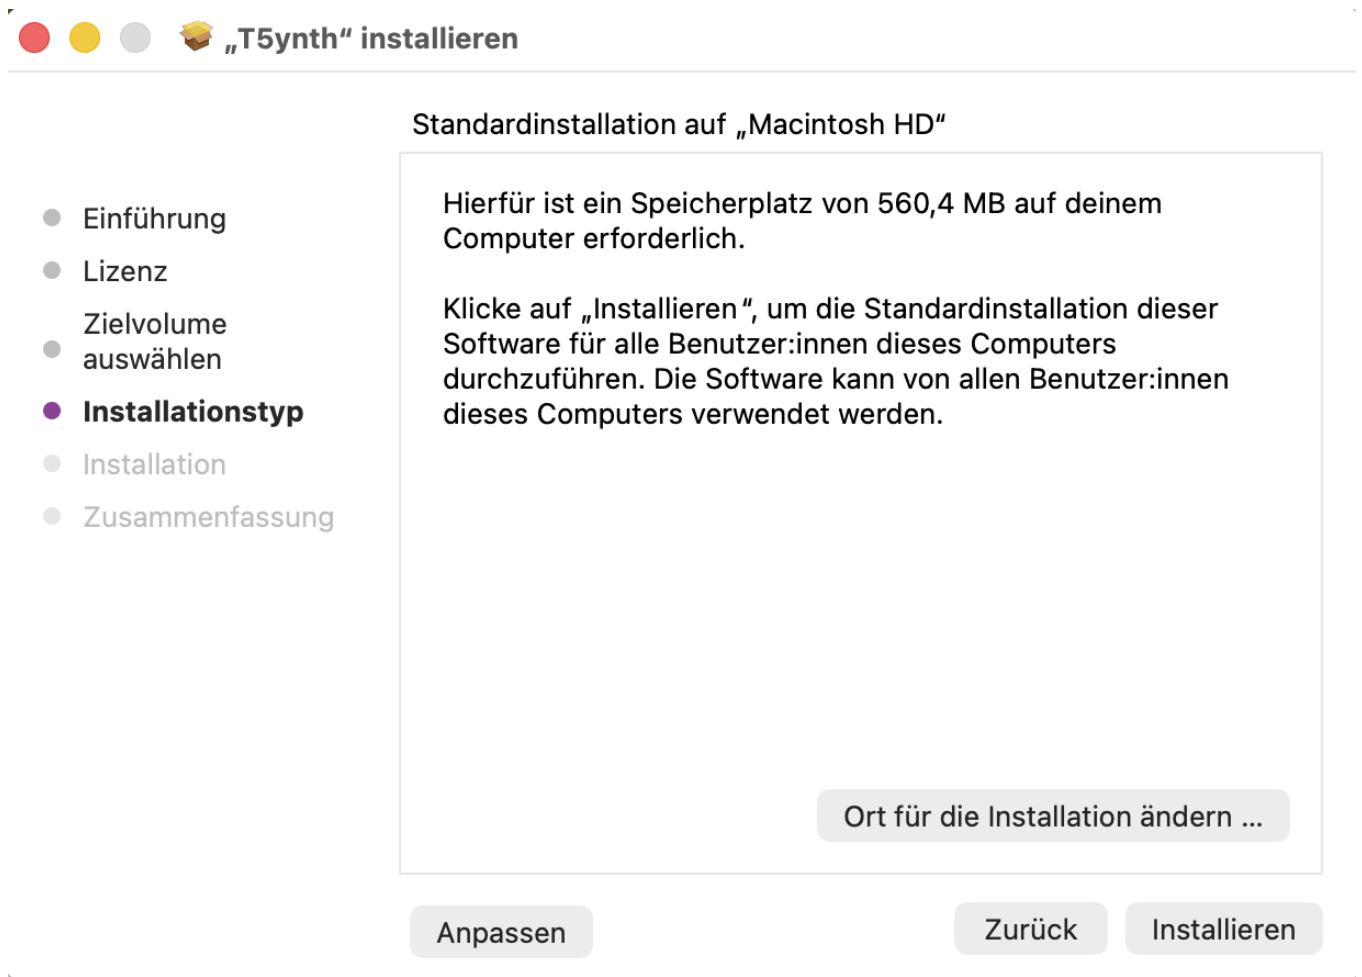
Step 7: Choose the install target

In the standard case, choose installation for all users of this computer.



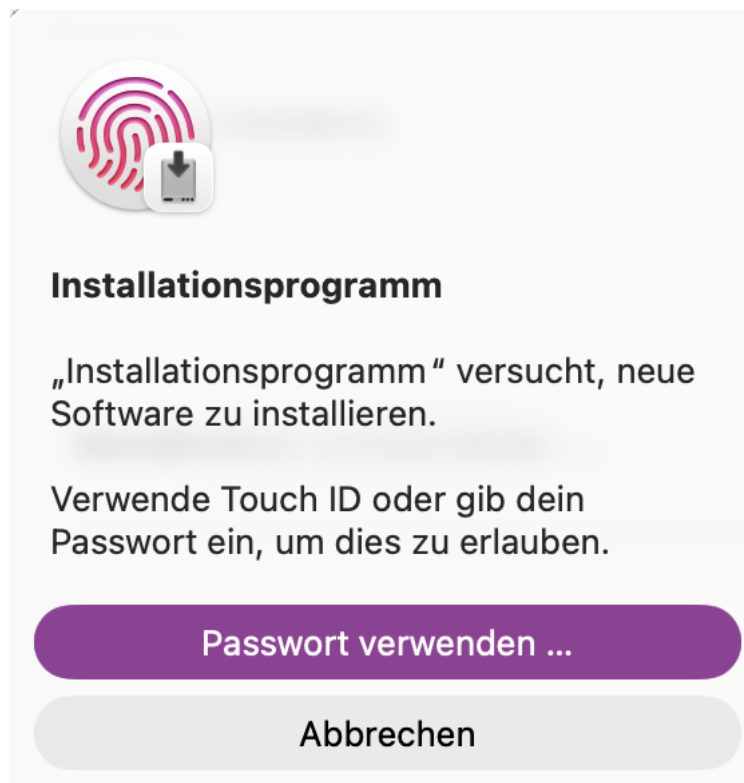
Step 8: Start the installation

Click "Installieren" to begin the standard installation on the system volume.



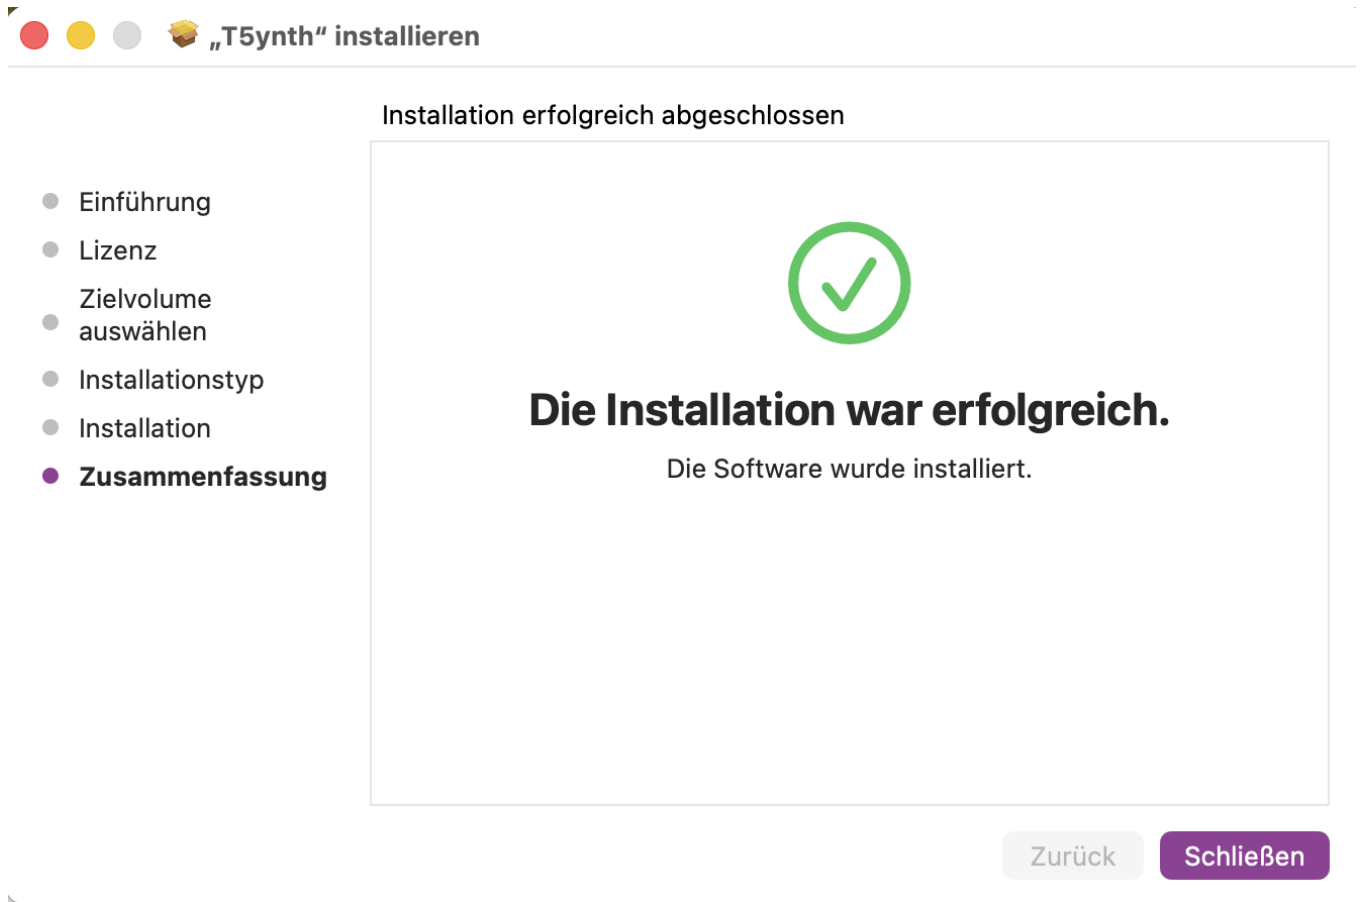
Step 9: Confirm with password or Touch ID

macOS asks for administrator permission for the actual installation.



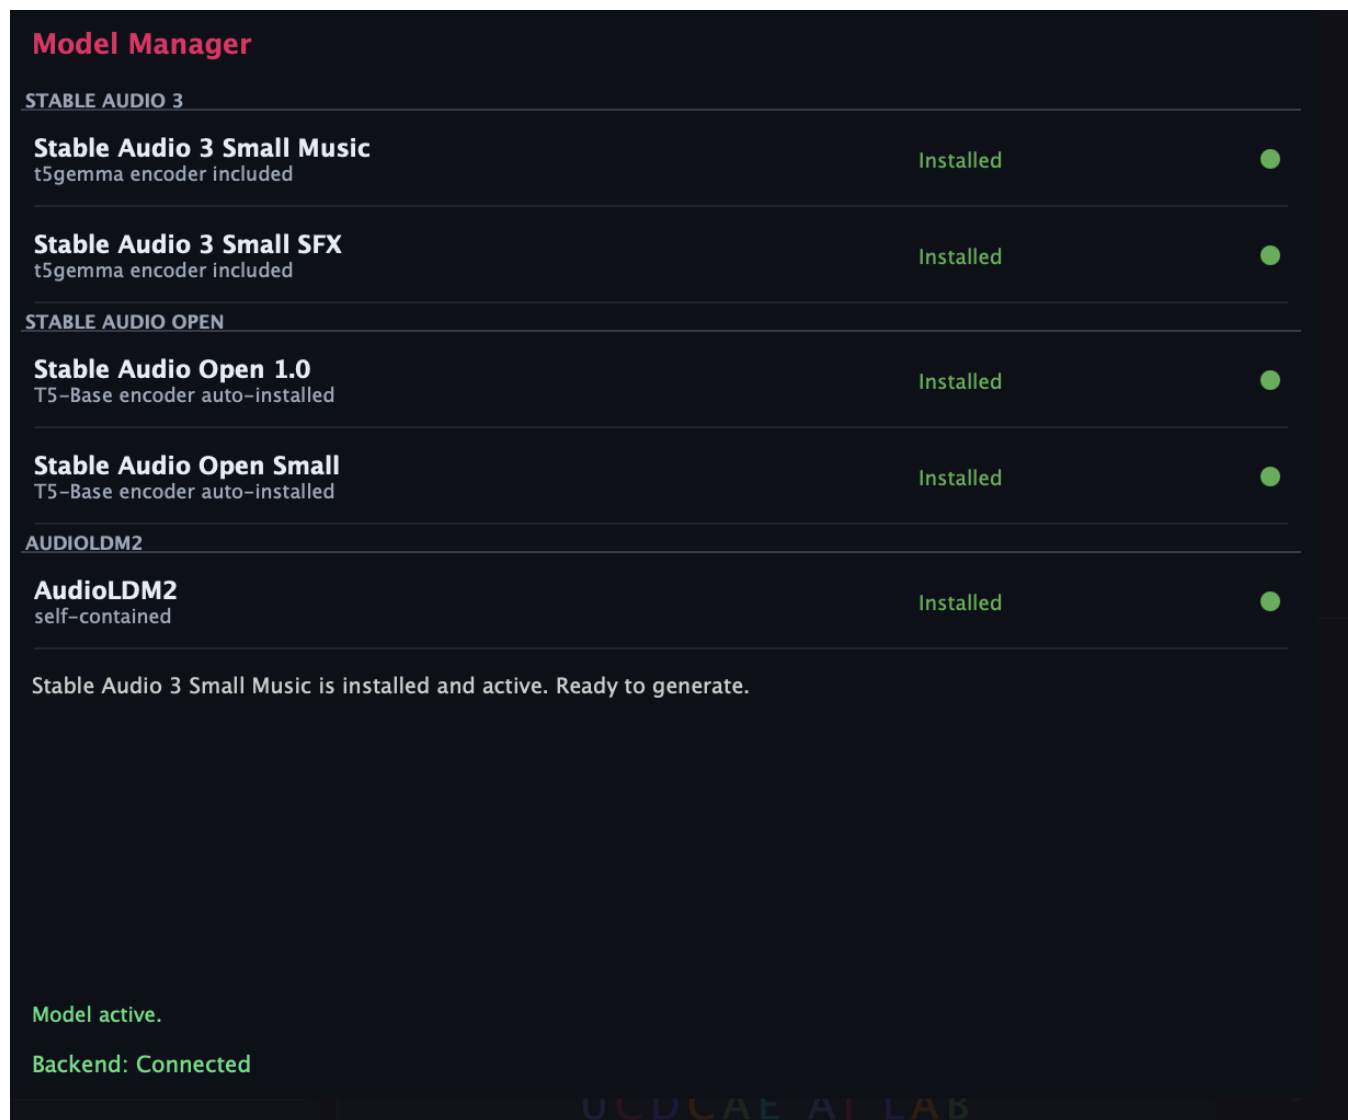
Step 10: Installation complete

Once this dialog appears, T5ynth.app is correctly installed in /Applications. Then launch T5ynth and download a model in the Model Manager — every engine installs directly in the app (see the appendices).



The Model Manager: download models in the app

All models are downloaded directly in the app: Settings > Model Manager. Each row shows an engine with its name, the required text encoder (installed automatically with it), a status text and a light on the right – green means fully installed. To install, select the engine, accept the license with “Accept & Download” and wait for the progress bar; the light then turns green, and “Backend: Connected” at the bottom means it is ready to use. This screenshot shows the target state with all five engines installed.



Appendix C: Loading Stable Audio 3 (recommended default)

Stable Audio 3 is T5ynth's current default model (version 2.1.0). There are two variants: Small Music for instrumental music and Small SFX for sound effects; both share one architecture and one t5gemma text encoder.

Before downloading, T5ynth shows two licenses that both must be accepted: the Stability AI Community License for the audio model and the Google Gemma Terms of Use plus Prohibited Use Policy for the t5gemma encoder. The weights then download without a HuggingFace account or token.

- Select “Stable Audio 3 Small Music” or “Small SFX” in the Model Manager.
- Accept both licenses with “Accept & Download”; the download starts.
- Both variants share the same t5gemma text encoder.
- Status switches to “Installed”, then “Backend: Connected”.

Note

Stable Audio 3 reads up to 256 tokens per prompt. T5ynth automatically prepends the required modality prefix (“TrackType: Music, ...” or “TrackType: SFX, ...”). SA3 was trained on comma-joined, field-tagged metadata, so structured prompts such as “Instruments: synth pad, Moods: warm” fit its training distribution; plain descriptions work too.

Appendix B: Loading Stable Audio Open 1.0

Stable Audio Open 1.0 is the original full-size Stable Audio Open checkpoint. It installs directly inside T5ynth, like Stable Audio 3.

Before downloading, T5ynth shows the Stability AI Community License, which must be accepted. The ~4.85 GB checkpoint then downloads without a HuggingFace account or token, and a copy of the license is written into the model folder.

- Select “Stable Audio Open 1.0” in the Model Manager.
- Accept the license with “Accept & Download”; the download starts.
- Status switches to “Installed”, then “Backend: Connected”.
- Already have the files from Stability by hand? Use “Auto-Scan” instead.

Note

Stable Audio Open 1.0 needs the “T5-Base text encoder” (ungated, Apache-2.0). T5ynth installs it automatically together with the model – no extra step. Stable Audio Open 1.0 is also much larger than Small; for first tests on a new Mac, Stable Audio 3 or Stable Audio Open Small is usually the quicker starting point.

Appendix D: Loading Stable Audio Open Small

Stable Audio Open Small is the compact, fastest Stable Audio Open checkpoint. It installs directly inside T5ynth, like Stable Audio 3 and Stable Audio Open 1.0 – no HuggingFace account and no manual download required any more.

Before downloading, T5ynth shows the Stability AI Community License, which must be accepted. The ~1.68 GB checkpoint then downloads without a HuggingFace account or token – from an ungated mirror that redistributes the official weights unmodified – and a copy of the license is written into the model folder.

- Select “Stable Audio Open Small” in the Model Manager.
- Accept the license with “Accept & Download”; the download starts.
- The required “T5-Base text encoder” (ungated, Apache-2.0) is installed automatically in the same step.
- Status switches to “Installed”, then “Backend: Connected”.
- Already have the files from Stability by hand? Use “Auto-Scan” instead.

Note

Stable Audio Open Small is the fastest model and already converges at around 8 steps; higher step counts bring no benefit here. CFG should stay at 1.

Appendix A: Loading AudioLDM2

AudioLDM2 is an academic latent-diffusion text-to-audio model (CVSSP / University of Surrey, Liu et al. 2023) for general audio, music and speech synthesis. It is self-contained and needs no separate text encoder.

Before downloading, T5ynth shows the license, which must be accepted. The model then downloads ungated directly from HuggingFace – no account and no token.

- Select “AudioLDM2” in the Model Manager.
- Accept the license with “Accept & Download”; the download starts.
- Status switches to “Installed”, then “Backend: Connected”.

Note

License: CC BY-NC-SA 4.0 – non-commercial use only, no revenue threshold, no exceptions. The recommended default model remains Stable Audio 3 (Appendix C).